


## Lecture 8 – MPI (Continued)

### Send and Receive



- Cooperative data transfer
- To (from) whom is data sent (received)?
- What is sent?
- How does the receiver identify it?

1

## Message Passing: Send

**MPI\_Send(address, count, datatype, dest, tag, comm)**

- (address, count) = a contiguous area in memory containing the message to be sent
- datatype = Type of data, e.g. integer, double precision (note that MPI had standard datatypes)
- dest = integer identifier representing the processor to send the message to
- tag = non-negative integer that the destination can use to selectively screen messages
- comm = communicator = group of processors

2

## Message Passing: Receive

**MPI\_Recv(address, count, datatype, source, tag, comm, status)**

- (address, count) = a contiguous area in message reserved for the message to be received
- datatype = Type of data, e.g. integer, double precision (note that MPI had standard datatypes)
- source = integer identifier representing the processor that sent the message
- tag = non-negative integer that the destination can use to selectively screen messages
- comm = communicator = group of processors
- status = information about the message that is received

3

## Single Program Multiple Data (SPMD)

- Proc 0 and Proc 1 are actually performing different operations
- However, not necessary to write separate programs for each processor
- Typically, use conditional statement and proc id to define the job of each processor:

```
integer :: a(10)

if(my_id == 0) then
    MPI_Send(a,10,MPI_INT,1,0,MPI_COMM_WORLD)
else if(my_id == 1) then
    MPI_Recv(a,10,MPI_INT,0,0,MPI_COMM_WORLD)
end if
```

4

## Different Types of Sends and Receives

- **There are 4 types of Sends and Receives:**
  - Standard (blocking)
  - Synchronous
  - Buffered
  - Ready
- **Each of these types can be performed in:**
  - Blocking mode
  - Non-blocking mode
- **The code programmer must make the choice of which type and mode to use depending on the circumstance**

5


## Send/Receive Types

- **Standard:** similar to Blocking except receive will not allow processor to continue only until its buffer can be reused.
- **Blocking:** receive will not allow processor to continue until it has received its message. Receive acts as a Barrier to that processor.
- **Synchronous:** send (or receive) does not start until a matching receive (or send) is posted indicating it is ready. Send acts as “blocking” until matching receive occurs. In this case, send acts as a Barrier for those processors.
- **Buffered:** either a system or user-defined buffer is made available for send/receive so that communication can proceed.

6

## Deadlock

- **Example: exchange data between 2 procs:**



- **MPI\_Send is a synchronous operation. If no system buffering is used, it keeps waiting until a matching receive is posted.**

7

## Deadlock

- **Both processors are waiting for each other → deadlock**
- **However, OK if system buffering exists → unsafe programming, however**
- **Note: MPI\_Recv is blocking and non-buffered**
- **Another real deadlock:**

Proc 0	Proc 1
MPI_Recv	MPI_Recv
MPI_Send	MPI_Send

- **Fix by reordering communication**

Proc 0	Proc 1
MPI_Send	MPI_Recv
MPI_Recv	MPI_Send

8

## Buffered/Nonbuffered Communications

- **No-buffering (phone calls)**
  - Proc 0 initiates the send request and rings Proc 1. It waits until Proc 1 is ready to receive. The transmission starts.
  - Synchronous communication – completed only when the message was received by the receiving proc
- **Buffering (beeper)**
  - The message to be sent (by Proc 0) is copied to a system-controlled block of memory (buffer)
  - Proc 0 can continue executing the rest of its program
  - When Proc 1 is ready to receive the message, the system copies the buffered message to Proc 1
  - Asynchronous communication – may be completed even though the receiving proc has not received the message

9

## Buffered Communication

- Buffering requires system resources, e.g. memory, and can be slower if the receiving proc is ready at the time of requesting the send
- **Application buffer:** address space that holds the data in the user's computer program
- **System buffer:** system space for storing messages. In buffered communication, data in application buffer is copied to/from system buffer
- **MPI allows communication in buffered mode:**  
MPI\_Bsend, MPI\_Ibsend
- **User allocates the buffer by:**  
MPI\_Buffer\_attach(buffer, buffer\_size)
- **Free the buffer by MPI\_Buffer\_detach**
- **An alternate to MPI\_Buffer commands is to allocate memory for buffer with standard allocate statement<sup>10</sup>**

## Blocking / Non-blocking Communication

- **Blocking Communication (McDonald's)**
  - The receiving proc has to wait if the message is not ready and has not received initial signal from sending proc
  - Different from synchronous communication (where sending proc will not begin sending until it has received explicit permission from receiving proc)
  - Proc 0 may have already buffered the message to system and Proc 1 is ready, but the interconnection network is busy
- **Non-blocking Communication (In & Out)**
  - Proc 1 checks with the system if the message has arrived yet. If not, it continues doing other stuff. Otherwise, get the message from the system.
- **Useful when computation and communication can be performed at the same time**
- **MPI allows both non-blocking send and receive**

11


## MPI\_Isend and MPI\_Irecv

- In non-blocking send, program identifies an area in memory to serve as a send buffer. Processing continues immediately without waiting for message to be copied out from the application buffer
- The user's program *should not* modify the application buffer until the non-blocking send has completed
- Non-blocking communication can be combined with non-buffering: MPI\_Issend, or buffering: MPI\_Ibsend
- Use MPI\_Wait or MPI\_Test to determine if the non-blocking send or receive has completed

12

## Example: Data Exchange in a Ring Topology

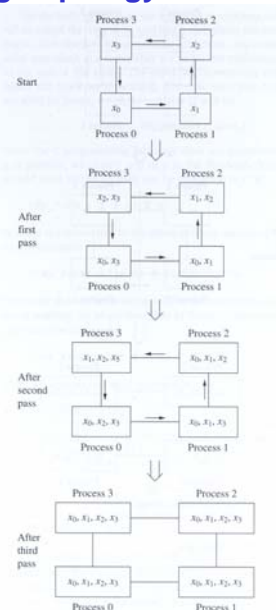
Analogous to MPI\_Allgather



- **Blocking version:**

```
for (i=0; i<p; i++) {
    send_offset = ((my_id-i+p)%p)*blksize;
    recv_offset = ((my_id-i-1+p)%p)*blksize;
    MPI_Send(y+send_offset,blksize,MPI_FLOAT,
            my_id+1,0,ring_com);
    MPI_Recv(y+recv_offset,blksize,MPI_FLOAT,
            my_id-1,0,ring_com,&status);
}
```

From Pacheco

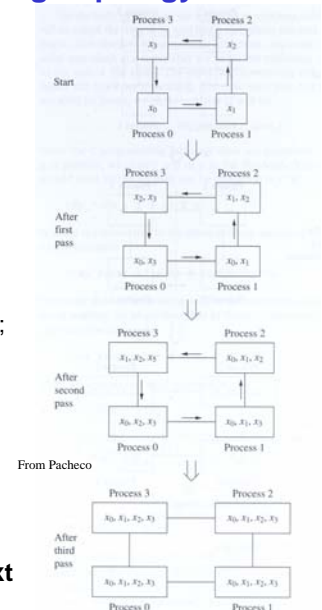



## Example: Data Exchange in a Ring Topology

- **Non-Blocking version:**

```
send_offset = my_id*blksize;
recv_offset = (my_id-1+p)*blksize;
for (i=0; i<p; i++) {
    MPI_Isend(y+send_offset,blksize,MPI_FLOAT,
            my_id+1,0,ring_com,&send_request);
    MPI_Irecv(y+recv_offset,blksize,MPI_FLOAT,
            my_id-1,0,ring_com,&recv_request);
    send_offset = ((my_id-i-1+p)%p)*blksize;
    recv_offset = ((my_id-i-2+p)%p)*blksize;
    MPI_Wait(&send_request,&status);
    MPI_Wait(&recv_request,&status);
}
```

- The communication and computations of next offsets are overlapped.



## Summary of Communication Modes

- **4 communication modes in MPI: standard (blocking), buffered, synchronous, ready.** They can be either blocking or non-blocking
- In standard (blocking) modes (MPI\_Send, MPI\_Recv,...), it is up to the system to decide whether messages should be buffered. Note there is a limited, finite amount of memory for system buffers.
- In synchronous mode, a send will not complete until a matching receive has been posted which has begun reception of the data
  - MPI\_Ssend (blocking), MPI\_ISsend (non-blocking)
  - No system buffering
- In buffered mode, the completion of a send does not depend on the existence of a matching receive
  - MPI\_Bsend (blocking), MPI\_IBsend (non-blocking)
  - System buffering by MPI\_Buffer\_attach and MPI\_Buffer\_detach
- Ready mode not discussed


15

## Collective Communications

- Communication pattern involving all the procs; usually more than 2
- **MPI\_Barrier:** synchronize all processors
- **Broadcast (MPI\_Bcast)**
  - A single proc sends the same data to every other proc
- **Reduction (gather/add) (MPI\_Reduce)**
  - All the procs contribute data that is combined using a binary operation
  - Example: max, min, sum, etc.
  - One proc obtains the final answer
- **Allreduce (MPI\_Allreduce)**
  - Same as MPI\_Reduce but every proc contains the final answer
  - Effectively as MPI\_Reduce + MPI\_Bcast, but more efficient

16

## An Implementation of the “Max” Function



- Tree-structured communication: (find the maximum among procs)
- Only needs  $\log_2 p$  stages of communication
- Not necessarily optimum on a particular architecture

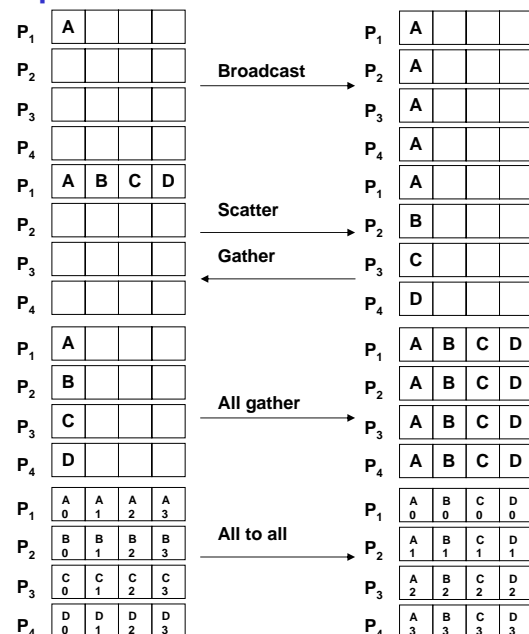
17

## Other Collective Communicators

- **Scatter (MPI\_Scatter)**
  - Split the data on the root processor into p segments
  - The 1<sup>st</sup> segment is sent to proc 0, the 2<sup>nd</sup> to proc 1, etc.
  - Similar to but more general than MPI\_Bcast
- **Gather (MPI\_Gather)**
  - Collect the data from each processor and store the data on root processor
  - Similar to but more general than MPI\_Reduce
- **Can collect and store the data on all procs using MPI\_Allgather**

18

## Comparison of Collective Communicators



19

## Homework 3

- Finish reading Chap. 1-3 of Using MPI by Gropp et al.
- Look at the parallel routine to compute  $\pi$  (calcpip) in the Codes directory (and the Examples directory on Wopr)
- Due Friday Oct. 23: Modify the calcpip.f routine to do a more accurate integration (Simpson's Rule) described in the next slide
  - Provide a listing of all subroutines and test this algorithm for different numbers of processors using parallel run.qsub batch submit procedure on wopr
  - Provide the CPU time as a function of the number of processors (up to 8) and the answer

20

### Homework 3

A more accurate alternative to the trapezoidal rule is Simpson's rule. The basic idea is to approximate the graph of  $f(x)$  by arcs of parabolas rather than line segments. Suppose that  $p < q$  are real numbers, and let  $r$  be the midpoint of the segment  $[p, q]$ . If we let  $h = (q - p)/2$ , then an equation for the parabola passing through the points  $(p, f(p))$ ,  $(r, f(r))$ , and  $(q, f(q))$  is

$$y = \frac{f(p)}{2h^2}(x - r)(x - q) - \frac{f(r)}{h^2}(x - p)(x - q) + \frac{f(q)}{2h^2}(x - p)(x - r).$$

If we integrate this from  $p$  to  $q$ , we get

$$\frac{h}{3}[f(p) + 4f(r) + f(q)].$$

Thus, if we use the same notation that we used in our discussion of the trapezoidal rule and we assume that  $n$ , the number of subintervals of  $[a, b]$ , is even, we can approximate

$$\int_a^b f(x)dx \doteq \frac{h}{3}[f(x_0) + 4f(x_1) + 2f(x_2) + 4f(x_3) + \cdots + 2f(x_{n-2}) + 4f(x_{n-1}) + f(x_n)].$$

Assuming that  $n/p$  is even, write

- a serial program and
- a parallel program that uses Simpson's rule to estimate  $\int_a^b f(x)dx$ .

From Parallel  
Programming  
with MPI by  
Pacheco