

Lecture 11 – Compilers

- **We want to make the compiler's job easy:**

1. Overhead:

Subroutine calls, indirect memory references, test within loops, wordy tests, variables preserved unnecessarily

2. Things that restrict compiler analysis:

Pointers, subroutine calls, indirect memory references

A good algorithm may give you far more speed than tuning!!

1

Numerical Libraries

- **Use the libraries provided by the vendors.**

Several smart people spent weeks (months) tuning the implementation and rewriting critical sections in assembler.

- **There are now some self-tuning libraries (FFTW, ATLAS,...) that have performances comparable to hand-tuned code.**

For more information, see:

http://en.wikipedia.org/wiki/Automatically_Tuned_Linear_Algebra_Software

2

Efficient Programming

- **Design cache-friendly data structures**

- Make each cache-line transfer pay
- Nested loops - Fortran leftmost subscript varies fastest
- Nested loops - C rightmost subscript varies fastest

- **Avoid aliasing – arguments pointing to same memory**

- **Ensure loops can be analyzed**

- Avoid branches within innermost loop
- choose "counted" loops (which include DO or IF loops, but not DO WHILE loops)

3

Blocking

Transpose: $A = B^T$


```
DO I=1,N
  DO J=1,N
    A(J,I)=B(I,J)
  END DO
END DO
```

A accessed with unit stride,
B with non-unit stride N



```
DO J=1,N
  DO I=1,N
    A(J,I)=B(I,J)
  END DO
END DO
```

B accessed with unit stride,
A with non-unit stride N



We want to block references, processing small blocks of A and B, to conserve cache entries.



4

Possible Parallel Programming Models

- **Shared memory:**

- Automatic parallelization
- Pthreads (Coarse, Medium Grain)
- Compiler directives (Coarse, Medium Grain): OpenMP for cores including new Intel Mic (Many Integrated Core processor) and even GPUs, OpenACC for GPUs
- Graphical processors (Fine Grain) : CUDA (C or Fortran), OpenCL (C)

- **Message passing:**

- MPI: message passing interface
- PVM: parallel virtual machine
- HPF: high performance fortran

5

Shared-Memory Pthreads

- **Posix threads:**

- Standard definition but non-standard implementation
- Hard to code
- Set of low-level primitives that give user more control at the price of complication
- Predominantly used by system programmers rather than application programmers

A website where you can learn more about Posix threads is:
<http://www.lynuxworks.com/products/posix/threads.php3>

6

Shared-Memory OpenMP

- **New “de-facto” standard available on all major platforms**
- **Easy to implement (easier than distributed memory parallelization)**
 - Higher-level primitives in the form of compiler directives
 - Easier to use for application programmers
- **Single source for parallel and serial code**
- **OpenMP is included in most Fortran90/95 compilers**
 - Is included in pgf90 on wopr
 - Latest versions can also be used for GPUs

A website where you can learn more about OpenMP is:
<http://www.openmp.org>

7

Shared Memory on GPUs (CUDA, OpenCL, OpenACC)

- **Graphical processors (GPUs) have their own shared memory**
- **Processing on the GPUs can be performed with standards that include**
 - CUDA (in C or Fortran) for NVIDIA GPUs
 - OpenCL for any GPU (NVIDIA or other vendors)
 - Both languages use the concept of Threads similar to OpenMP which now can be used directly for GPUs as well as the OpenACC directive package
- **Algorithm Kernels are called from a “kernel” routine that calls another “thread” routine**

8

Shared Memory on GPUs (CUDA, OpenCL, OpenACC)

- Thread routines consist of instructions that are performed for every micro-processor

Websites where you can learn more about CUDA or OpenCL are

<http://developer.nvidia.com/opengl>

<http://developer.nvidia.com/category/zone/cuda-zone>

9

Distributed-Memory MPI

- **Standard parallel interface:**
 - MPI 1: two-sided communication
 - MPI 2: extend MPI 1 with single side communication (GAS)
 - See <http://www.cs.berkeley.edu/~bonachea/upc/mpl2.html> for a good explanation
- **Need to rewrite the code to perform message passing**
- **Code portable on all architectures (open-source code available)**
- **There are different implementations of MPI depending on the computer: OpenMPI, MPICH, Portland Group MPI (with pgi)**

Websites where you can find out more about OpenMPI and MPICH (open-source standards) is:

<http://www.open-mpi.org/>

<http://www-unix.mcs.anl.gov/mpi/mpich>

10

Distributed-Memory PVM

- **Parallel Virtual Machine**
- **Another popular message passing interface**
- **Useful in environment with multiple vendors and for MPMD approach**
- **Again, PVM is open-source software that can be downloaded**

The website where you can find out more about PVM is:

http://www.csm.ornl.gov/pvm/pvm_home.html

11

Simple Code to Compute π

$$\int_0^1 \frac{4}{1+x^2} dx = \pi$$

```
program compute_pi
integer n, i
double precision w, x, sum, pi, f, a
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
print *, 'Enter number of intervals: '
read *, n
! calculate the interval size
w = 1.0d0/real(n)
sum = 0.0d0
do i = 1, n
    x = w * (real(i) - 0.5d0)
    sum = sum + f(x)
end do
pi = w * sum
print *, 'computed pi = ', pi
stop
end
```

12

Pthreads C-code 1/3

```
#include <stdio.h>
#include <unistd.h>
#include <sys/times.h>
#include <pthread.h>
#include <stdlib.h>

float sumall, width;
int i, iend;
__private int ibegin;
__private int cut;
__private float x;
__private float xsum;
void *do_work();
main(argc,argv)
int argc;
char *argv[];
{
    /* Pi - Program loops over slices in interval, summing area of each slice */
    struct tms time_buf;
    float ticks, t1, t2;
    int intrvl, numthreads, istart;
    pthread_t pt[32];

    /* get intervals from command line */

    intrvl = atoi(argv[1]);
    numthreads = atoi(getenv( "PL_NUM_THREADS"));
    printf(" intervals = %d PL_NUM_THREADS = %d\n",intrvl, numthreads);
```

13

Pthreads C-code 2/3

```
/* get number of clock ticks per second and initialize timer */

    ticks = sysconf(_SC_CLK_TCK);
    t2 =times(&time_buf);

    /* -- Compute width of cuts */

    width = 1. / intrvl;
    sumall = 0.0;

    /* -- Loop over interval, summing areas */
    istart = 1;
    iend = intrvl/numthreads;
    for (i = 0; i < numthreads - 1 ; i++)
    {
        pthread_create(&pt[i], pthread_attr_default, do_work,(void *) istart);
        istart += iend;
    }
    do_work( istart);
    istart += iend;
    for (i = 0; i < numthreads - 1 ; i++)
    {
        pthread_join(pt[i], NULL);
    }

    /* -- finish any remaining slices */
    iend = intrvl - (intrvl/numthreads)* numthreads;
    if( iend) do_work( istart);
    /* -- Finish overall timing and write results */
    t1 = times(&time_buf);
    printf("Time in main = %20.14e sum = %20.14f\n", (t1 -t2)/ticks,sumall);
    printf("Error = %20.15e\n",sumall - 3.14159265358979323846);
}
```

14

Pthreads C-code 3/3

```
void *do_work(istart)
int istart;
{
    ibegin = istart;
    xsum = 0.0;
    for (cut = ibegin ; cut < ibegin+iend; cut++)
    {
        x = ((float) cut) - .5) * width ;
        xsum += width * 4. /(1. + x * x);
    }
    sumall += xsum;
}
```

- As you can see, Posix (pthreads) gives a lot of control at the expense of complexity. This is probably not the technique to use in engineering simulations.
- OpenMP is much simpler!!

15

OpenMP code 1/1

```
program compute_pi
integer n, i
double precision w, x, sum, pi, f, a
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
print *, 'Enter number of intervals: '
read *,n
! calculate the interval size
w = 1.0d0/real(n)
sum = 0.0d0
!$OMP PARALLEL DO PRIVATE(x), SHARED(w)
!$OMP REDUCTION(+: sum)
do i = 1, n
    x = w * (real(i) - 0.5d0)
    sum = sum + f(x)
end do
!$OMP END PARALLEL DO
pi = w * sum
print *, 'computed pi = ', pi
stop
end
```

- Directives are much simpler !!

16

MPI code 1/2

```

program compute_pi
include 'mpif.h'
double precision mypi, pi, w, sum, x, f, a
integer n, myid, numprocs, i, rc
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
call MPI_INIT( ierr )
call MPI_COMM_RANK( MPI_COMM_WORLD, myid, ierr )
call MPI_COMM_SIZE( MPI_COMM_WORLD, numprocs, ierr )
if ( myid .eq. 0 ) then
    print *, 'Enter number of intervals: '
    read *, n
endif
call MPI_BCAST(n,1,MPI_INTEGER,0,MPI_COMM_WORLD,ierr)
! calculate the interval size
w = 1.0d0/real(n)
sum = 0.0d0
do i = myid+1, n, numprocs
    x = w * (real(i) - 0.5d0)
    sum = sum + f(x)
enddo
mypi = w * sum

```

17

MPI code 2/2

```

! collect all the partial sums
call
MPI_REDUCE(mypi,pi,1,MPI_DOUBLE_PRECISION,MPI_SUM,0,
$ MPI_COMM_WORLD,ierr)
! node 0 prints the answer.
if (myid .eq. 0) then
    print *, 'computed pi = ', pi
endif
call MPI_FINALIZE(rc)
stop
end

```

- We will use MPI across the CPU/Cores in this course since it is relatively simple, readily available for many computers with both Fortran 95/03 and C, and has become the standard.

- However, we could instead use OpenMP within a node (blocks).

- OR, we could use a combination of MPI across the nodes and OpenMP within the node (blocks).

18

PVM master code 1/3

```

program compute_pi_master
include '~/pvm3/include/tpvm3.h'
parameter (NTASKS = 5)
parameter (INTERVALS = 1000)
integer mytid
integer tids(NTASKS)
real sum, area
real width
integer i, numt, msgtype, bufid, bytes, who, info

sum = 0.0
! Enroll in PVM
call pvmfmytid(mytid)
! spawn off NTASKS workers
call pvmfspawn('compmpi.worker', PVMDEFAULT, '', NTASKS, tids, numt)
width = 0.0
i = 0
! Multi-cast initial dummy message to workers
msgtype = 0
call pvmfinit(0, info)
call pvmfpack(INTEGER4, i, 1, 1, info)
call pvmfpack(REAL4, width, 1, 1, info)
call pvmfmcast(NTASKS, tids, msgtype, info)
! compute interval width
width = 1.0 / INTERVALS

```

19

PVM master code 2/3

```

do i = 1, NTASKS
! for each interval, 1) receive area from worker 2) add area to sum 3) send worker new interval
number and width
call pvmfrecv(-1, -1, bufid)
call pvmfbuinfo(bufid, bytes, msgtype, who, info)
call pvmfunpack(REAL4, area, 1, 1, info)
sum = sum + area
call pvmfinit(0, info)
call pvmfpack(INTEGER4, i, 1, 1, info)
call pvmfpack(REAL4, width, 1, 1, info)
call pvmfpack(REAL4, area, 1, 1, info)
call pvmfpack(REAL4, width, 1, 1, info)
call pvmfpack(REAL4, area, 1, 1, info)
enddo

! Signal to workers that tasks are done
i = -1
call pvmfinit(0, info)
call pvmfpack(INTEGER4, i, 1, 1, info)
call pvmfpack(REAL4, width, 1, 1, info)

```

20

PVM master code _{3/3}

```
! Collect the last NTASK areas and send the completion signal
do i = 1, NTASKS
  call pvmfrecv(-1,-1, bufid)
  call pvmfbuinfo(bufid, bytes, msgtype, who, info)
  call pvmfunpack(REAL4, area, 1, 1, info)

  sum = sum + area

  call pvmfsend(who, msgtype, info)
enddo

print 10,sum
10 format(x,'Computed value of Pi is ',F8.6)

call pvmfexit(info)
end
```

Note that the PVM code is generally more complicated than it's MPI counter-part due to buffer manipulation

21

PVM worker code _{1/2}

```
integer mytid, master
real area
real width, int_val, height
integer int_num
! Enroll in PVM
call pvmfmytid(mytid)
! who is sending me work?
call pvmfparent(master)

! receive first job from the master
call pvmfrecv(-1, -1, info)
call pvmfunpack(INTEGER4, int_num, 1, 1, info)
call pvmfunpack(REAL4, width, 1, 1, info)

! While I've not been sent the signal to quit, I'll keep processing
40 if (int_num .eq. -1) goto 50
! compute interval value from interval number
int_val = int_num * width
! compute height of given rectangle
height = F(int_val)

! compute area
area = height * width
```

22

PVM worker code _{2/2}

```
! send area back to master
call pvmfinit send(PvmDataDefault, info)
call pvmfpack(REAL4, area, 1, 1, info)
call pvmfsend(master, 9, info)

! Wait for next job from master
call pvmfrecv(-1, -1, info)
call pvmfunpack(INTEGER4, int_num, 1, 1, info)
call pvmfunpack(REAL4, width, 1, 1, info)
goto 40
! all done
50 call pvmfexit(info)
end

real function f(x)
real x
f = 4.0/(1.0+x**2)
return
end
```

This is one reason why MPI has now become the standard for distributed-memory parallel

23