

Lecture 16a – Parallel/Processor Domain Decomposition

- Parallelizing a computer code requires the following steps:

- Identifying portions of the work that can be performed concurrently
 - Mapping the concurrent pieces of work into multiple processes (that run in parallel)
 - Distributing the input, output, and intermediate data associated with the program
 - Managing accesses to data shared by multiple processors
 - Synchronizing the processors at various stages of the parallel program execution
- Parallel/Processor Domain Decomposition
- MPI, PVM, or OpenMP

1

Granularity

- Granularity:** The number and size of tasks into which a problem can be decomposed

- A decomposition into a large number of small tasks is considered *fine-grained*
- A decomposition into a small number of large tasks is considered *coarse-grained*

- Degree of Granularity:**

- The maximum number of tasks that can be executed simultaneously in a parallel program at any given time is the *maximum degree of concurrency*
- The average number of tasks that can concurrently run over the entire duration of program execution is the *average degree of concurrency*

2


Critical Path

- We can create a **task-dependency graph** of the various tasks performed during program execution


- A task-dependency graph is an abstraction (flow chart) used to show dependencies among program tasks and their relative order of execution
 - The initial tasks in the graph are called the *start nodes*
 - The final tasks in the graph are called the *finish nodes*
 - The longest path between any pair of start and finish nodes is the *critical path*
 - The sum of the weights (amount of computational work) of the nodes (tasks) along the critical path is known as the *critical path length*
 - The ratio of the total amount of work to the critical path length is the *average degree of concurrency*

3

Examples



- Critical Path = 2 (4-6-7) or (1-9-8)
- Critical Path Length = 27 (1-9-8)
- Total Amount of Work = 63
- Average Degree on Concurrency = $63/27 = 2.33$



- Critical Path = 3 (4-5-6-7)
- Critical Path Length = 34
- Total Amount of Work = 64
- Average Degree on Concurrency = $64/34 = 1.88$

4

Granularity

- At first, it appears that the time required to solve a problem can be reduced by increasing the granularity of decomposition and perform more tasks in parallel.
- **This is not the case in general, however.**
- There is an inherent bound on how fine-grained a decomposition of a problem permits (due to the nature of the problem, for example: the number of intervals, number of data points, number of cells, number of blocks, etc.)
- Another major factor that limits speedup from further decomposition is the **interaction** between tasks (processors)

5

Interaction

- The pattern of interaction among the tasks can be described in a **task-interaction graph**
 - The nodes in a task-interaction graph are the tasks
 - The edges connecting the nodes represent the interactions. The direction of the edge-lines represent the flow of data (from one task to another).
 - The nodes and edges can be assigned weights based upon the amount of work performed.
- The task-interaction graph looks much the same as the task-dependency graph but with weights assigned to the interactions (edges).

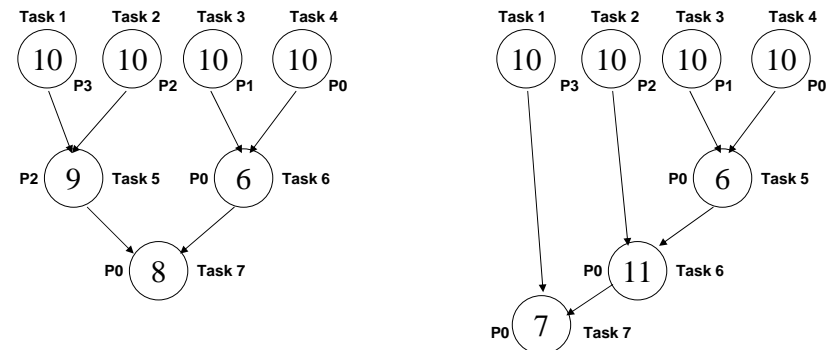
6

Mapping

- A **process** is a computing agent that performs tasks (much like a processor)
- The procedure by which tasks are assigned to processes for execution is called **mapping**
- A good mapping should
 - Maximize the use of concurrency by mapping independent tasks onto different processes
 - Minimize the total completion time by ensuring that processes are available to execute tasks on the critical path as soon as possible
 - Minimize the interaction among processes by mapping tasks with a high degree of mutual interaction onto the same process

7

Examples: Mapping 7 Tasks onto 4 Processes



- A maximum of 4 processes can be used efficiently. The maximum degree of concurrency is 4
- The last 3 tasks can be mapped arbitrarily to satisfy the constraints of the task-dependency graph
- Makes more sense to map the tasks connected by an interaction edge onto the same process since this prevents an inter-task interaction from becoming an inter-processes interaction

8

Decomposition Techniques

- **There are different techniques for decomposition depending on the type of problem that is being solved**
 - Recursive decomposition
 - Data-decomposition
 - Exploratory decomposition
 - Speculative decomposition
- **Recursive and data-decomposition are the most general**
- **Exploratory and Speculative are for special-purpose problems**

9

Recursive Decomposition

- **Recursive decomposition used for problems that can be solved with a divide-and-conquer strategy**
 - Example: Sorting
 - First, divide the problem into a set of independent sub-problems
 - Each of these sub-problems is solved by sub-dividing again by another number of sub-problems, and so on...

10

Data Decomposition

- **Data decomposition is the most general way to break a problem into sub-problems**
 - The data on which the computations are performed is partitioned
 - The data partitioning is used to induce a partitioning of the computations in the tasks
 - For instance, break the data (intervals, cells, nodes, blocks, etc.) into partitions and perform the tasks of a computation on the partitions of the data. This is the technique typically used in simulations and engineering computations and will be what you use in Projects-4 and 5.
 - Data decomposition can be performed based upon:
 - Output data
 - Input data
 - Output and input data

11

Exploratory Decomposition

- **Exploratory decomposition is used for problems whose underlying computations correspond to a search of a space for solutions.**
 - Partition the search space into smaller parts
 - Search each one of these parts concurrently until the desired solutions are found
 - Optimization or search problems could fall into this category

12

Speculative Decomposition

- **Speculative decomposition is used when a program may take on one of many possible computational branches depending on the output of other computations that precede it.**
 - While one task is performing the computation whose output is used in deciding the next computation, other tasks can concurrently start the computations of the next stage
 - Example: switches (branches) of logic
 - While one task is performing the computation that will eventually resolve the switch, other tasks could pick up the multiple branches of the switch in parallel
 - When the input for the switch has finally been computed, the computation corresponding to the correct branch would be used while that corresponding to the other branches would be discarded