

## Lecture 20 – Multi-Processor Performance

- CPU scalability
- Memory scalability
- Interconnection network
- Bandwidth and latency issues
- Problem size and granularity
- *How many processors can we use efficiently?*
- In the last class, we discussed Amdahl's law for single and multiple processors. Let's first go back and review the multiple-processor model.

1

## Amdahl's Law - Parallel Processing

- Ideally, if a computation can be carried out in  $p$  equal parts, the total execution time will be nearly  $1/p$  of the time required by a single processor
- Suppose  $t_j$  denotes the wall clock time required to execute a task with  $j$  processors
- Speedup,  $S_p$ , for  $p$  processors is defined as

$$S_p = \frac{t_1}{t_p}$$

- Where  $t_1$  is the time required for the most-efficient sequential algorithm to complete the calculation, and  $t_p$  is the time required for the most efficient parallel implementation of the same algorithm, from beginning to end, using  $p$  processors.

2

## Amdahl's Law - Parallel Processing

- The *computational efficiency* using  $p$  processors is defined as

$$E_p = \frac{S_p}{p} \quad 0 \leq E_p \leq 1$$

- Then, the *total execution time* using  $p$  processors is given by

$$t_p = \frac{ft_1}{p} + (1-f)t_1 = \frac{t_1(f + (1-f)p)}{p} \geq (1-f)t_1$$

3

## Amdahl's Law - Parallel Processing

- Therefore, the speedup on  $p$  processors is then

$$S_p = \frac{p}{(f + (1-f)p)} \leq \frac{1}{1-f}$$

(Ware's Law)

- This shows that the speedup is considerable reduced even for pretty large values of  $f$  (close to 95%)

4

## Amdahl's Law - Parallel Processing


- **Parallel overhead** is the additional amount of work that is required on the parallel implementation of a sequential algorithm arising from the use of a parallel computer:
  - Inter-processor communication
  - Load imbalance
  - Additional computation resulting from the algorithm that is being parallelized not being as efficient as the most efficient serial algorithm.
- If the total time spent in solving a problem over all processing elements is  $pT_p$  and  $T_s$  is the time spent doing useful work (consider this the time for a single processor), then  $T_o$  is the overhead time:

$$T_o = pT_p - T_s$$

- This is the overhead time spent in communication

5

## Amdahl's Law - Parallel Processing



6

## Amdahl's Law - Parallel Processing

- Amdahl's law is a simplistic, yet powerful way of looking at the problem of **scalability**.
- In a naïve way, it points out that a large number of processors cannot be used on any computational task, since  $f$  needs to be very close to 1. For example,  $f=0.999$  would allow the use of a maximum number of processors equal to 1000.


7

## Example - Adding on a Hypercube

- Consider adding  $n$  numbers using  $n$  processors of a hypercube ( $n$  is a power of two)
- Initially, each processor gets one number, and at the end, one processor has the sum of all of them
- Addition and communication take each 1 unit of time.
  - The addition can be performed in some constant time,  $t_c$ .
  - The communication of a single word can be performed in time  $t_s + t_w$  where
    - $t_s$  is the start-up time to handle message at the sending and receiving nodes
    - $t_w$  is the per-word transfer time

8

### Example - Adding on a Hypercube (sum 16 numbers on 16 processors)



### Example - Adding on a Hypercube

$$T_p = \Theta(\log n) \quad S_p = \Theta\left(\frac{n}{\log n}\right) \quad E_p = \frac{S_p}{n} = \Theta\left(\frac{1}{\log n}\right)$$

- **Performance not impressive:**

- For 16 numbers and 16 processors,  $E_p = 25\%$  (Note that all log functions in Introduction to Parallel Computing by Grama et al are in base 2. See Appendix A)

$$\log_2 x = \frac{\log_{10} x}{\log_{10} 2} = \frac{\log_{10} x}{.301}$$

- **Maybe we are using too many processors for this problem? Maybe the problem size is too small?**

- **Let's try again.**

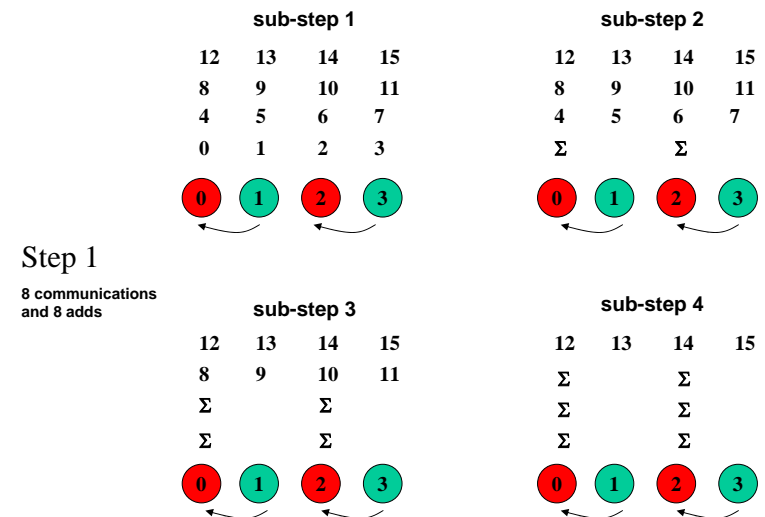
10

### Example - Adding on a Hypercube

- Assign larger pieces of data to each processor (more than one number per processor). We are increasing the work-load of each processor by doing this.
- Consider adding  $n$  numbers using  $p$  processors of a hypercube,  $p < n$ , both  $p$  and  $n$  are powers of two. Using fewer processors than the maximum number of elements is called *scaling down*.
- Initially, each processor gets  $n/p$  numbers, and at the end, one processor has the sum of all of them
- Addition and communication take each 1 unit of time


11

### Example - Adding on a Hypercube (Non-optimal Algorithm) (Sum 16 numbers on 4 processors – mimic operations of 16 processors)




12

### Example - Adding on a Hypercube (Non-optimal Algorithm) (sum 16 numbers on 4 processors (non-optimally))



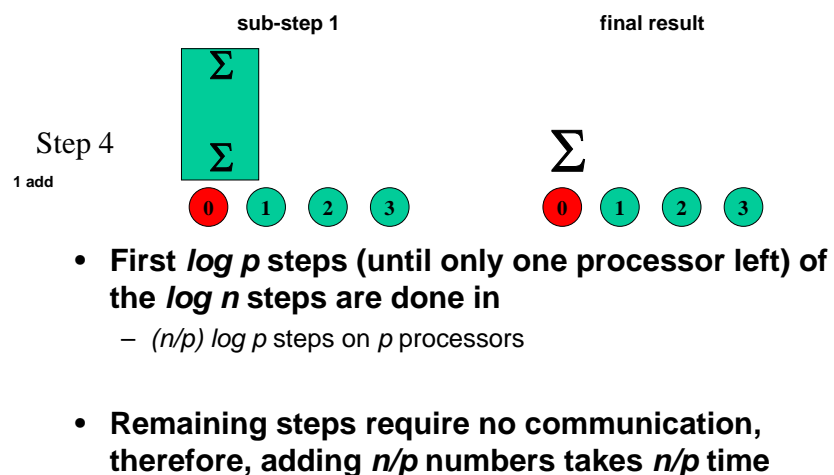
13

### Example - Adding on a Hypercube (Non-optimal Algorithm) (sum 16 numbers on 4 processors (non-optimally))



14

### Example - Adding on a Hypercube (Non-optimal Algorithm) (sum 16 numbers on 4 processors (non-optimally))



12 communications  
and 15 adds

15

### Example - Adding on a Hypercube (sum 16 numbers on 4 processors (non-optimally))

- Expected computational time from our previous example of  $n$ -processor hypercube:

$$T_p = \Theta\left(\frac{n}{p} \log n\right)$$

- However what we find in this analysis for a  $p$ -processor hypercube and  $n$ -words:

$$T_p = \Theta\left(\frac{n}{p} \log p\right)$$

and the speedup and efficiency would be

$$S_p = \Theta\left(\frac{n}{\frac{n}{p} \log p}\right) = \Theta\left(\frac{p}{\log p}\right) \quad E_p = \frac{S_p}{n} = \Theta\left(\frac{p}{n \log p}\right)$$

- For 16 words on 4 processors,  $E_p = 12\%$  (worse – because it is non-optimal). Increasing work-load doesn't help!!

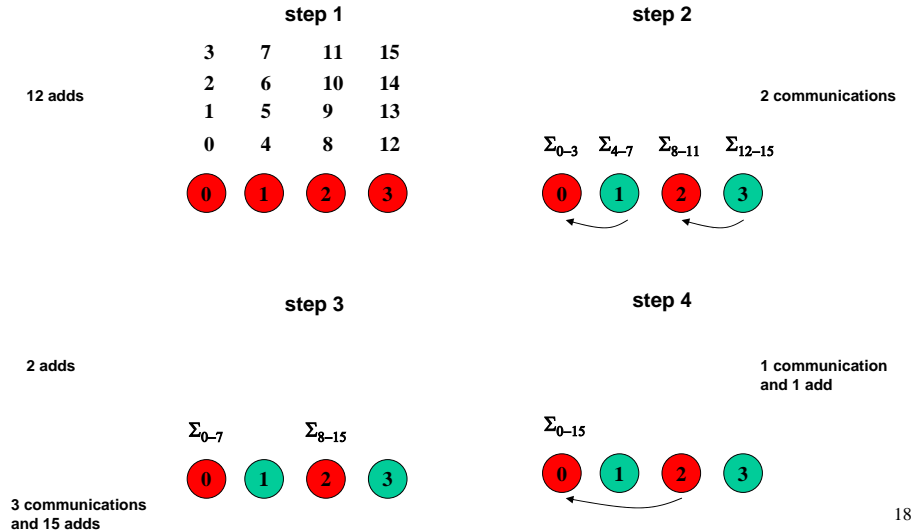
16

### Example - Adding on a Hypercube (sum 16 numbers on 4 processors (non-optimally))

- This example shows that if you have a bad parallel algorithm, increasing work-load will not necessarily help
- We must go back and change the algorithm to make it more optimal

17

### Example - Adding on a Hypercube (Optimal) (Sum 16 numbers on 4 processors optimally)



18

### Example - Adding on a Hypercube (sum 16 numbers on 4 processors optimally)

- Total computation time, assuming it takes one unit of time to add two numbers and one unit of time to transfer a number between neighboring processors

$$T_p = \left( \frac{n}{p} - 1 \right) + 2 \log p \approx \frac{n}{p} + 2 \log p$$

- Speedup

$$S_p = \frac{n}{\frac{n}{p} + 2 \log p} = \frac{np}{n + 2p \log p}$$

19

### Example - Adding on a Hypercube (sum 16 numbers on 4 processors optimally)

- Efficiency:  $E_p = \frac{S_p}{p} = \frac{n}{n + 2p \log p}$

- For 16 numbers on 4 processors,  $E_p = 50\%$  (much better!!!)

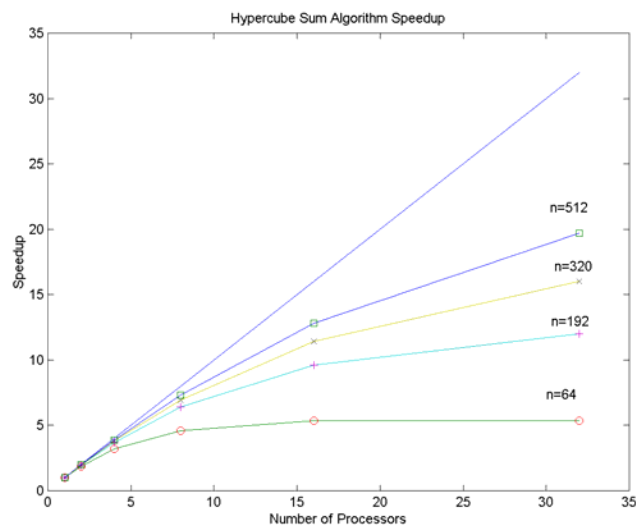
P	Sp	Ep
2	1.6	0.8
4	2.0	0.5
8	2.0	0.25
16	1.78	0.11

- Note that:

- Speedup does not increase linearly with the number of processors
- Speedup increases with larger vector sizes (for fixed  $p$ )
- If we scale the problem at the same time we increase the number of processors, and efficiency remains constant, we have a *scalable system and algorithm*

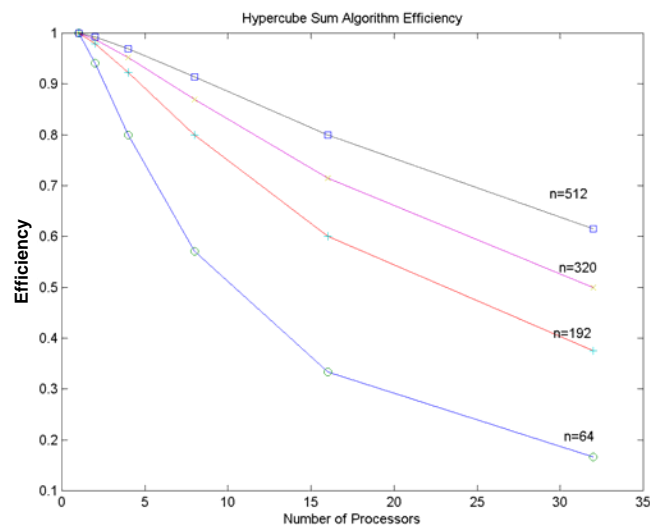
20

## Example - Adding on a Hypercube



21


## Example - Adding on a Hypercube



22


## Realistic Test Case

- A multi-block flow solver. Not very different!!!



23

## A Scalable Version of Another Multi-Block Flow Solver



24

## Granularity

- Often, the concept of *granularity* is what matters most.
- **Granularity**: Ratio of amount of *communication* that a processor requires to the amount of *computation* that it performs
- **Coarse grain** parallelism (low granularity): typical MPI applications. Decrease communication overhead. Most suitable for low performance networks (but also for high performance, of course).

25

## Granularity

- **Fine grain** parallelism (high granularity): most typical in OpenMP and GPU applications where small portions of computation (such as operations inside a loop) are divided among multiple processors.
- Fine grain parallelism requires high performance network/memory subsystems.

26

## Granularity

- Note the inherent benefit in scaling up the problem size for a block of typical dimension,  $L$ :
  - Amount of *communication* is proportional to the *surface area* of the block  $\approx L^2$
  - Amount of *computation* is proportional to the *volume* of the block  $\approx L^3$
- Therefore, as the problem gets larger, the granularity decreases  $\approx \frac{1}{L}$

27

## Bandwidth

- **Bandwidth** is the *rate of information transfer* that a communication subsystem can maintain (Mb/sec)
  - Bandwidth that matters is *software bandwidth* (using MPI, for example)
  - More is better (... and much more expensive)
  - To minimize communication overhead, send only necessary information
  - Typical number: 60 Mb/sec for Origin2000, 85 Mb/sec for Matrx, 775 Mb/sec for Davistron, 440 Mb/sec for Vortex or Wopr

28

## Latency

- **Latency** is defined as the time it takes to send a zero-length message from one processor to another (measured in microseconds, typically)
  - Less is better (... and more expensive)
  - Software latency (MPI, for example) matters
  - Impacts short messages mostly (and algorithms that rely heavily on short messages, i.e. multigrid)
  - Try to agglomerate messages if possible to decrease communication cost.
  - Typical number: 15 microseconds for SGI Origin2000, 50 microseconds for Matrx, 1.3 microseconds for Davistron, 2 microseconds for Vortex

29

## Homework 7

- Read Chapter 7 of Introduction to Parallel Computing by Grama et al
  - Read about shared-memory parallel computing and OpenMP
  - A good tutorial and users' manual on OpenMP is:  
<https://computing.llnl.gov/tutorials/openMP/>

30