

## Lecture 21 – Shared-Memory Parallel Computing

- **Now let's start to discuss how we can use shared-memory systems to perform parallel computing**
- **Remember that in shared-memory systems, all of the processors see all of the data**

1

## Motivation

- **Popularity of shared memory systems is increasing:**
  - Distributed Shared-Memory (DSM) computers (SGI Altix, SUN Ultra 10000, HP Exemplar) have replaced vector systems in many supercomputing centers
- **Compiler directed parallelism is attracting attention:**
  - Single version of sequential and parallel code
  - Industrial standard is now available: OpenMP

2

## Motivation

- **New generation of parallel machines with multiple CPUs per node**
  - IBM SP3 (older technology):
    - each node is a 16-way shared memory machine
    - there is a single bus to the high performance network per node
  - IBM BlueGene/L (later technology):
    - Each node is a 2-way shared memory machine
    - 1024 nodes per rack, 16-128 I/O nodes per rack, 3D Torus interconnect 350 Mb/s bandwidth, 1.5  $\mu$ s latency
  - Quad-, Octo-, and now higher Core Chip-Sets (eg. Mic, latest technology)
- **Mixed programming model:**
  - OpenMP in each node
  - MPI across nodes

3

## Programming Models

- **Shared memory options:**
  - Automatic parallelization (some compilers)
  - Pthreads (POSIX threads)
  - Compiler directives: **OpenMP**
- **Message passing options:**
  - **MPI**: message passing interface
  - PVM: parallel virtual machine
  - HPF: high performance Fortran

4

## OpenMP

- **OpenMP is an API for writing multithreaded applications:**
  - A set of compiler directives and library routines for parallel application programmers
  - Makes it easy to create multithreaded programs in Fortran, C and C++

5

## OpenMP Supporters

- **Hardware vendors:**
  - Compaq/HP, IBM, Intel, SGI, SUN,...
- **Software vendors**
  - KAI (now Intel), PGI, Absoft, PSR,...

<http://openmp.org/wp>

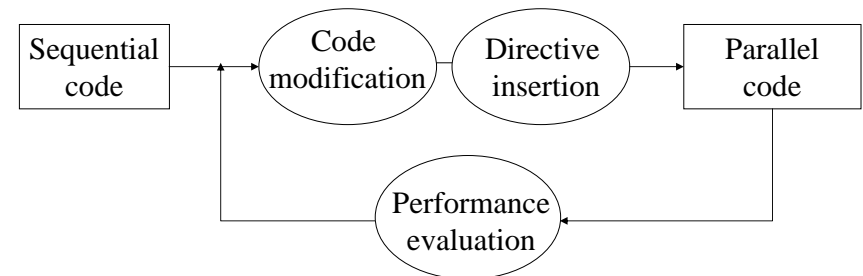
6

## OpenMP

- Fine grained parallelism (at loop level)
- Coarse grained parallelism
- Compiler directives, library and environment variables *extend* base language
- NOT automatic parallelization
- Since the constructs are directives, an OpenMP program can be compiled by compilers that do not support OpenMP

7

## Incremental Parallelization



Work can be done incrementally!!!!!!

8

## Serial Code

```

program compute_pi
integer n, i
double precision w, x, sum, pi, f, a
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
print *, 'Enter number of intervals: '
read *, n
! calculate the interval size
w = 1.0d0/n
sum = 0.0d0
do i = 1, n
    x = w * (i - 0.5d0)
    sum = sum + f(x)
end do
pi = w * sum
print *, 'computed pi = ', pi
stop
end

```

9

## MPI Code

```

program compute_pi
include 'mpif.h'
double precision mypi, pi, w, sum, x, f, a
integer n, myid, numprocs, i, rc
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
call MPI_INIT( ierr )
call MPI_COMM_RANK(MPI_COMM_WORLD,
&               myid, ierr )
call MPI_COMM_SIZE(MPI_COMM_WORLD,
&               numprocs, ierr )
if ( myid .eq. 0 ) then
    print *, 'Enter number of intervals: '
    read *, n
endif
call MPI_BCAST(n,1,MPI_INTEGER,0,
&             MPI_COMM_WORLD,ierr)

! calculate the interval size
w = 1.0d0/n
sum = 0.0d0
do i = myid+1, n, numprocs
    x = w * (i - 0.5d0)
    sum = sum + f(x)
enddo
mypi = w * sum
! collect all the partial sums
call MPI_REDUCE(mypi,pi,1,
&               MPI_DOUBLE_PRECISION,
&               MPI_SUM,0,
&               MPI_COMM_WORLD,ierr)
! node 0 prints the answer
if (myid .eq. 0) then
    print *, 'computed pi = ', pi
endif
call MPI_FINALIZE(rc)
stop
end

```

10

## OpenMP Code

```


program compute_pi
integer n, i
double precision w, x, sum, pi, f, a
! function to integrate
f(a) = 4.d0 / (1.d0 + a*a)
print *, 'Enter number of intervals: '
read *, n
! calculate the interval size
w = 1.0d0/n
sum = 0.0d0
!$OMP PARALLEL DO PRIVATE(x), SHARED(w,n)
!$OMP& REDUCTION(+: sum)
do i = 1, n
    x = w * (i - 0.5d0)
    sum = sum + f(x)
end do
!$OMP END PARALLEL DO
pi = w * sum
print *, 'computed pi = ', pi
stop
end

```

11


## Execution Model

- **Master thread spawns a team of threads as needed**
  - Concept of threads also used in GPU parallel computing
- **Parallelism is added incrementally**



12

## Execution Model



- OpenMP programs start as single thread
- In parallel regions, additional threads are created
- Master thread is part of the team
- Outside parallel regions, additional threads go away (or sleep)

13

## Directives or pragmas

- **Fortran (fixed format)**

!\$omp ...

c\$omp ...

\*\$omp ...

If the character in column 6 is different from a space or a 0 it is a continuation line

- **Fortran (free form)**

– A line that begins with !\$omp is an OpenMP directive

– A directive that needs to be continued on the next line:

!\$OMP ..... &

- **C:**

#pragma omp

14

## Conditional Compilation

- The selective disabling of OpenMP construct applies only to directives.
- Application may contain statements that are specific to OpenMP but not intended for OpenMP

### Fortran :

- A line beginning with a *sentinel* (!\$,c\$,\*\$ in fixed format, !\$ in free format) is ignored if not in OpenMP

!\$ iam=omp\_get\_num\_thread()

### C:

- Preprocessor macro \_OPENMP
- #pragma omp

15

## Programming Model

- **OpenMP is a shared memory model**

- Threads communicate by sharing variables

- **Unintended sharing of data can lead to race conditions**

- Race conditions: the program's outcome changes as the threads are scheduled differently

16

## Example in C

### Sequential code

```
void main()
{
    double a[100000];
    for (int i=0,i<100000,i++){
        do_something(a[i]);
    }
}
```

### Parallel code

```
void main()
{
    double a[100000];
    #pragma omp parallel for
    for (int i=0,i<100000,i++){
        do_something(a[i]);
    }
}
```

17

## Example in Fortran

### Sequential code

```
real*8 a(100000)
do i = 0,100000
    do_something(a)
enddo
```

### Parallel code

```
real*8 a(100000)
!#omp parallel do
do i = 0,100000
    do_something(a)
enddo
```

18

## OpenMP Constructs

- **OpenMP's constructs (directives) fall into 5 categories:**
  - Parallel regions
  - Worksharing
  - Synchronization
  - Data environment
  - Runtime functions/environment variables
- **We will begin to discuss these constructs in the next lecture**

19